

Class

BQOM 2578 | Data Mining

Theresa Wohlever

2025-01-01

Table of contents

Executive Summary	1
Loading packages	1
Importing data	2
Preliminary Analysis	2
Logistic Regression Models	4
Confusion Matrix	7
Comparing Models	18
Train and Test datasets	21

Executive Summary

Loading packages

Lets start by calling some libraries that are useful for Logistic Regressions.

- caTools for splitting the data in a smart way.
- ROCR for creating ROC and AUC curves.
- caret for Machine Learning in general.

Importing data

We are using a *loans* dataset from Taiwan, *Credit_data.csv*, which is available on Canvas. Download it to your working folder and use `getwd()` and `setwd()` if you need to change your working directory.

Here is information for some of the key variables:

- Limit: Amount of given credit, in 1,000 New Taiwan Dollars (1 USD is approx. 30 NTD, checked in Sept, 2025)
- Gender: 1 = male; 2 = female;
- Marital status: 1 = married; 2 = single/other;
- Late i: whether the person was late i months ago;
- Default: dependent variable (0/1); if the client defaulted (1) or not (0)

Note how gender and Marital Status have 1 and 2 rather than 0 and 1 as we prefer for dummies. We can let R handle this for us.

```
table(loans$Default)
```

```
  0    1  
23045 6610
```

```
summary(loans$Default)
```

```
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   
 0.0000  0.0000  0.0000  0.2229  0.0000  1.0000
```

```
paste("The proportion of customers defaulting is: ",round(mean(loans$Default),2))
```

```
[1] "The proportion of customers defaulting is:  0.22"
```


Preliminary Analysis

Let's begin with evaluating the Correlation Matrix

```
loans<-loans%>%relocate(Default) # moves the variable "Default" to the first column (left  
head(loans)
```

	Default	Limit	Gender	MaritalStatus	Age	Late1	Late2	Late3
1	1	20	2		1 24	1	1	0
2	1	120	2		2 26	0	1	0
3	0	90	2		2 34	0	0	0
4	0	50	2		1 37	0	0	0
5	0	50	1		1 57	0	0	0
6	0	50	1		2 37	0	0	0

```
cormat <- round(cor(loans),2)
corrplot(cormat, method="number", type="upper",diag=FALSE, tl.cex=1.2)
```



Logistic Regression Models

Formula is very simple: `glm(y ~ X, family="binomial")`.

For variables that we want to treat as factors (categorical variables) we use `as.factor`. R will change it to a dummy, taking the lowest value as 0.

Specifically, after using `as.factor(Gender)`, 1 will become 0 and 2 will become 1.

```
logreg0 = glm(Default ~ Limit, data=loans, family="binomial")
summary(logreg0)
```

Call:

```
glm(formula = Default ~ Limit, family = "binomial", data = loans)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.7458015	0.0224716	-33.19	<2e-16 ***
Limit	-0.0033003	0.0001265	-26.09	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 31467 on 29654 degrees of freedom
 Residual deviance: 30696 on 29653 degrees of freedom
 AIC: 30700

Number of Fisher Scoring iterations: 4

```
logreg = glm(Default ~ Limit + as.factor(Gender) + as.factor(MaritalStatus)
              + Age + Late1 + Late2 + Late3, data=loans, family="binomial")
summary(logreg)
```

Call:

```
glm(formula = Default ~ Limit + as.factor(Gender) + as.factor(MaritalStatus) +
    Age + Late1 + Late2 + Late3, family = "binomial", data = loans)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.5306414	0.0838129	-18.263	< 2e-16 ***
Limit	-0.0019926	0.0001346	-14.800	< 2e-16 ***
as.factor(Gender)2	-0.1372312	0.0313593	-4.376	1.21e-05 ***
as.factor(MaritalStatus)2	-0.1661281	0.0346078	-4.800	1.58e-06 ***
Age	0.0041690	0.0018393	2.267	0.0234 *
Late1	1.3578076	0.0412525	32.915	< 2e-16 ***
Late2	0.2976810	0.0550675	5.406	6.45e-08 ***
Late3	0.7161429	0.0472809	15.147	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 31467 on 29654 degrees of freedom
Residual deviance: 26975 on 29647 degrees of freedom
AIC: 26991

Number of Fisher Scoring iterations: 4

Let's examine the coefficients β and $\exp(\beta)$:

```
coef(logreg)
```

(Intercept)	Limit	as.factor(Gender)2
-1.530641432	-0.001992608	-0.137231191
as.factor(MaritalStatus)2	Age	Late1
-0.166128052	0.004168986	1.357807557
Late2	Late3	
0.297681026	0.716142876	

```
exp(coef(logreg))
```

(Intercept)	Limit	as.factor(Gender)2
0.2163968	0.9980094	0.8717687
as.factor(MaritalStatus)2	Age	Late1
0.8469378	1.0041777	3.8876605
Late2	Late3	
1.3467321	2.0465243	

```
coeftable<-data.frame(col1=coef(logreg),col2=exp(coef(logreg)))  
colnames(coeftable)<-c('Coefficient (log-odds)', 'e^coefficient (odds)')
```

```
coeftable
```

	Coefficient (log-odds)	e^coefficient (odds)
(Intercept)	-1.530641432	0.2163968
Limit	-0.001992608	0.9980094
as.factor(Gender)2	-0.137231191	0.8717687

as.factor(MaritalStatus)2	-0.166128052	0.8469378
Age	0.004168986	1.0041777
Late1	1.357807557	3.8876605
Late2	0.297681026	1.3467321
Late3	0.716142876	2.0465243

Confusion Matrix

What we get from the logistic regression are predicted probabilities. What we need to convert them into classification decisions is a threshold or cutoff.

```
loans$PredProbs1<-predict(logreg, newdata=loans, type="response")
# type="response" gives the probability, otherwise the output would be log odds.
summary(loans$PredProbs1)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.03034	0.11892	0.14713	0.22290	0.25709	0.74341

```
#We then transform that prediction into either 1 (True) or 0 (False) using a cutoff point
cutoff<-0.25
#Try different cutoff points:
#cutoff<-0.15 # Note the Probs don't change, just the classification
#cutoff<-0.05

loans$PredDefault1<-ifelse(loans$PredProbs1>=cutoff,1,0)
summary(loans$PredDefault1)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.0000	0.0000	0.0000	0.2547	1.0000	1.0000

```
paste("For a cutoff point of",cutoff, "the proportion of customers classified as defaulting
```

```
[1] "For a cutoff point of 0.25 the proportion of customers classified as defaulting is 0.25"
```

```
head(loans)
```

	Default	Limit	Gender	MaritalStatus	Age	Late1	Late2	Late3	PredProbs1
1	1	20	2		1 24	1	1	0	0.5119526
2	1	120	2		2 26	0	1	0	0.1588194
3	0	90	2		2 34	0	0	0	0.1333579
4	0	50	2		1 37	0	0	0	0.1661378
5	0	50	1		1 57	0	0	0	0.1989867
6	0	50	1		2 37	0	0	0	0.1621731

	PredDefault1
1	1
2	0
3	0
4	0
5	0
6	0

```
tail(loans)
```

	Default	Limit	Gender	MaritalStatus	Age	Late1	Late2	Late3	PredProbs1
29650	1	80	1		2 34	1	1	1	0.6586309
29651	0	220	1		1 39	0	0	0	0.1410708
29652	0	150	1		2 43	0	0	0	0.1398671
29653	1	30	1		2 37	1	1	1	0.6833773
29654	1	80	1		1 41	1	0	0	0.4597587
29655	1	50	1		1 46	0	0	0	0.1917780

	PredDefault1
29650	1
29651	0
29652	0
29653	1
29654	1
29655	0

Let's see what is the combination of Predicted vs Observed Default values (Try with different cutoffs above.)

```
paste("Observed Default values:")
```

```
[1] "Observed Default values:"
```



```
table(loans$Default)
```

```
      0      1  
23045 6610
```

```
paste("Predicted Default values:")
```

```
[1] "Predicted Default values:"
```

```
table(loans$PredDefault1)
```

```
      0      1  
22101 7554
```

```
paste("Observed by predicted values (confusion matrix) for cutoff of", cutoff, "is:")
```

```
[1] "Observed by predicted values (confusion matrix) for cutoff of 0.25 is:"
```

```
#first variable is rows, second variable is columns
```

```
ConfMatrix<-table(loans$Default,loans$PredDefault1) # table() with two variables will cross  
rownames(ConfMatrix)<-c("Obs False", "Obs True")  
colnames(ConfMatrix)<-c("Pred False", "Pred True")  
ConfMatrix
```

	Pred False	Pred True
Obs False	19198	3847
Obs True	2903	3707

The below function does that. Easier to copy and paste in the future. Use it to test multiple cutoff points.

```
#defining a new function called ConfMatrix with three inputs: Actual, Predicted and Cutoff
ConfMatrix_func<-function(actual_value,predicted_prob,cutoff){
  predicted_value<-ifelse(predicted_prob>=cutoff,1,0) # Classify by evaluating probability
  ConfMatrix<-table(actual_value,predicted_value)
  rownames(ConfMatrix)<-c("Obs False", "Obs True")
  colnames(ConfMatrix)<-c("Pred False", "Pred True")

  print(paste("For a cutoff of", cutoff,":"))
  print("Actual values in the test dataset")
  print(table(actual_value))
  print("Predicted values in the test dataset")
  print(table(predicted_value))

  return(ConfMatrix)
}
```

Let's try it out:

```
ConfMatrix_func(loans$Default,loans$PredProbs1,0.25)
```

```
[1] "For a cutoff of 0.25 :"
```

```
[1] "Actual values in the test dataset"
```

```
actual_value
```

```
  0    1
```

```
23045 6610
```

```
[1] "Predicted values in the test dataset"
```

```
predicted_value
```

```
  0    1
```

```
22101 7554
```

```

              predicted_value
actual_value Pred False Pred True
Obs False    19198    3847
Obs True     2903     3707
```

```
#ConfMatrix_func(loans$Default,loans$PredProbs1,0.15)
#ConfMatrix_func(loans$Default,loans$PredProbs1,0.05)
```

Now, calculate the performance measures of the confusion matrix:

```

predicted_value<-loans$PredDefault1
actual_value<-loans$Default

#True Negative means it was predicted negative and it is indeed negative.
TN<-sum(predicted_value==0 & actual_value==0)
#False Negative means it was predicted negative and it is actually positive.
FN<-sum(predicted_value==0 & actual_value==1)
#False Positive means it was predicted true and it is actually false.
FP<-sum(predicted_value==1 & actual_value==0)
#True Positive means it was predicted true and it is indeed true.
TP<-sum(predicted_value==1 & actual_value==1)

# Note that the logical operator "&" is used for an eval at the level of each element of the vector

print(paste("There are",TN," True Negatives"))

```

```
[1] "There are 19198  True Negatives"
```

```
print(paste("There are",FN," False Negatives"))
```

```
[1] "There are 2903  False Negatives"
```

```
print(paste("There are",FP," False Positives"))
```

```
[1] "There are 3847  False Positives"
```

```
print(paste("There are",TP," True Positives"))
```

```
[1] "There are 3707  True Positives"
```

```
print("As a result:")
```

```
[1] "As a result:"
```

```
#Getting Accuracy, Sensitivity, Specificity and Precision:
#Accuracy is the number of correct predictions over the total
ACCU<-(TN+TP)/(TN+TP+FN+FP)
#True Positive Rate (TPR), also called Sensitivity or recall is the number of correct posit
TPR<-TP/(TP+FN)
#True Negative Rate (TFR), also called Specificityis the number of correct negative predict
TNR<-TN/(TN+FP)
#Precision, also called Positive Predictive Value (PPV),is the number of TP over the total
PPV<-(TP)/(TP+FP)
print(paste("Accuracy is:", round(ACCU,4)))
```

```
[1] "Accuracy is: 0.7724"
```

```
print(paste("True Positive Rate is:", round(TPR,4)))
```

```
[1] "True Positive Rate is: 0.5608"
```

```
print(paste("True Negative Rate is:", round(TNR,4)))
```

```
[1] "True Negative Rate is: 0.8331"
```

```
print(paste("Precision PPV is:", round(PPV,4)))
```

```
[1] "Precision PPV is: 0.4907"
```

As you can imagine, there are packages that include functions to directly do what we did ‘manually’ coded above. In this case, the package is caret.

However, you need to have created the predicted values based on your selected cutoff PRIOR to calling this function.

```
# the PredDefault1 was calculated already with a chosen Cutoff
confusionMatrix(data=as.factor(loans$PredDefault1),reference=as.factor(loans$Default))
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	19198	2903

```

1  3847  3707

      Accuracy : 0.7724
      95% CI   : (0.7676, 0.7771)
No Information Rate : 0.7771
P-Value [Acc > NIR] : 0.9748

      Kappa : 0.3748

McNemar's Test P-Value : <2e-16

      Sensitivity : 0.8331
      Specificity : 0.5608
      Pos Pred Value : 0.8686
      Neg Pred Value : 0.4907
      Prevalence : 0.7771
      Detection Rate : 0.6474
      Detection Prevalence : 0.7453
      Balanced Accuracy : 0.6969

      'Positive' Class : 0

```

What?? These numbers are completely different! Why? Let's look at the help file (remember this is in the Caret Package)

```

# the table is flipped, AND the 0 and 1 are reversed too!!
# the parameter "positive" might be able to help us - let's try

confusionMatrix(data=as.factor(loans$PredDefault1),reference=as.factor(loans$Default), posi

```

Confusion Matrix and Statistics

```

      Reference
Prediction    0    1
      0 19198  2903
      1  3847  3707

      Accuracy : 0.7724
      95% CI   : (0.7676, 0.7771)
No Information Rate : 0.7771
P-Value [Acc > NIR] : 0.9748

```

```

Kappa : 0.3748

McNemar's Test P-Value : <2e-16

Sensitivity : 0.5608
Specificity : 0.8331
Pos Pred Value : 0.4907
Neg Pred Value : 0.8686
Prevalence : 0.2229
Detection Rate : 0.1250
Detection Prevalence : 0.2547
Balanced Accuracy : 0.6969

'Positive' Class : 1

```

That looks better!!

We can also just keep creating our own functions for our own purposes, in this case passing a Cutoff to the function call:

```

MyConfMatrixValues_func<-function(actual_value,predicted_prob,cutoff){

  predicted_value<-ifelse(predicted_prob>=cutoff,1,0)  # Apply the Cutoff

  ConfMatrix<-table(actual_value,predicted_value)
  rownames(ConfMatrix)<-c("Obs False", "Obs True")
  colnames(ConfMatrix)<-c("Pred False", "Pred True")

  print(paste("For a cutoff of", cutoff,":"))
  print("Actual values in the test dataset")
  print(table(actual_value))
  print("Predicted values in the test dataset")
  print(table(predicted_value))

  #predicted_value<-ifelse(predicted_prob>=cutoff,1,0)
  #True Negative means it was predicted negative and it is indeed negative.
  TN<-sum(predicted_value==0 & actual_value==0)
  #False Negative means it was predicted negative and it is actually positive.
  FN<-sum(predicted_value==0 & actual_value==1)
  #False Positive means it was predicted true and it is actually false.
  FP<-sum(predicted_value==1 & actual_value==0)

```

```

#True Positive means it was predicted true and it is indeed true.
TP<-sum(predicted_value==1 & actual_value==1)

print(paste("There are",TN," True Negatives"))
print(paste("There are",FN," False Negatives"))
print(paste("There are",FP," False Positives"))
print(paste("There are",TP," True Positives"))

#Getting Accuracy, Sensitivity, Specificity and Precision:
#Accuracy is the number of correct predictions over the total
ACCU<-(TN+TP)/(TN+TP+FN+FP)
#True Positive Rate (TPR), also called Sensitivity or recall is the number of correct posit
TPR<-TP/(TP+FN)
#True Negative Rate (TNR), also called Specificity is the number of correct negative predict
TNR<-TN/(TN+FP)
#Precision, also called Positive Predictive Value (PPV), is the number of TP over the total
PPV<-(TP)/(TP+FP)

print("As a result:")
print(paste("Accuracy is:", round(ACCU,4)))
print(paste("True Positive (Sensitivity) Rate is:", round(TPR,4)))
print(paste("True Negative (Specificity) Rate is:", round(TNR,4)))
print(paste("Precision PPV is:", round(PPV,4)))

  ConfMatrixValues<-data.frame(TN,FN,FP,TP,ACCU,TPR,TNR)
  return(ConfMatrix)
}

```

Now that we defined the function, let's use it:

```
MyConfMatrixValues_func(loans$Default, loans$PredProbs1, 0.25)
```

```

[1] "For a cutoff of 0.25 :"
```

[1] "Actual values in the test dataset"	
actual_value	
0	1
23045	6610

```

[1] "Predicted values in the test dataset"
```

predicted_value	
0	1
22101	7554

```

[1] "There are 19198 True Negatives"
[1] "There are 2903 False Negatives"
[1] "There are 3847 False Positives"
[1] "There are 3707 True Positives"
[1] "As a result:"
[1] "Accuracy is: 0.7724"
[1] "True Positive (Sensitivity) Rate is: 0.5608"
[1] "True Negative (Specificity) Rate is: 0.8331"
[1] "Precision PPV is: 0.4907"

```

		predicted_value	
actual_value		Pred False	Pred True
Obs False		19198	3847
Obs True		2903	3707

ROC (Receiver Operating Characteristic) Curve

Instead of manually trying different cutoff points, we can use the ROC Curve.

This also gives us the Area Under the Curve (AUC), which we can use to compare the model performance with other models:

```

# We'll use the function "prediction" (note: NOT predict),
# which transforms the predicted probabilities (first argument)
# and the actual 0/1 values (second argument)
# into a standardized format of class prediction, and store them into
# the object roc.pred...

roc.pred = prediction(loans$PredProbs1, loans$Default)
# ... which we can then use to actually create the ROC curve
# with the function "performance" (note: we need to store this
# so that we can then draw the curve):
perf = performance(roc.pred, "tpr", "fpr")
#If you don't get what it is doing: ?performance

plot(perf,                                     # the data
      main = "ROC Curve",                     # the chart's title
      xlab = "1 - Specificity",                # the name of the x-axis
      ylab = "Sensitivity",                    # the name of the y-axis
      colorize=TRUE)                           # add color to curve depending on threshold prob.

# ... and add the diagonal corresponding to the Random Assignment


```



```
# benchmark model:
```

```
abline(0,1) # adds line at intercept 0, with slope 1
```

ROC Curve



```
perf_auc = performance(roc.pred, "auc")  
as.numeric(perf_auc@y.values)
```

```
[1] 0.7454106
```

That is, our first model; logreg, has an area under the curve of 0.745. How does that compare to the baseline? (AUC=0.5).

How can we increase specificity? Let's say we want a specificity of at least 90%.

```
#That puts us in the green area of about 0.4
```

```
MyConfMatrixValues_func(loans$Default, loans$PredProbs1, 0.4)
```

```
[1] "For a cutoff of 0.4 :"
```

```
[1] "Actual values in the test dataset"
```

```
actual_value
```

```
0      1
```

```
23045 6610
```

```
[1] "Predicted values in the test dataset"
```

```
predicted_value
```

```

      0      1
24825 4830
[1] "There are 20955 True Negatives"
[1] "There are 3870 False Negatives"
[1] "There are 2090 False Positives"
[1] "There are 2740 True Positives"
[1] "As a result:"
[1] "Accuracy is: 0.799"
[1] "True Positive (Sensitivity) Rate is: 0.4145"
[1] "True Negative (Specificity) Rate is: 0.9093"
[1] "Precision PPV is: 0.5673"

```

	predicted_value	
actual_value	Pred False	Pred True
Obs False	20955	2090
Obs True	3870	2740

Comparing Models

Let's make another model with no demographic factors to avoid discrimination. Let's see the original and the new model.

```
summary(logreg)
```

Call:

```
glm(formula = Default ~ Limit + as.factor(Gender) + as.factor(MaritalStatus) +
    Age + Late1 + Late2 + Late3, family = "binomial", data = loans)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.5306414	0.0838129	-18.263	< 2e-16 ***
Limit	-0.0019926	0.0001346	-14.800	< 2e-16 ***
as.factor(Gender)2	-0.1372312	0.0313593	-4.376	1.21e-05 ***
as.factor(MaritalStatus)2	-0.1661281	0.0346078	-4.800	1.58e-06 ***
Age	0.0041690	0.0018393	2.267	0.0234 *
Late1	1.3578076	0.0412525	32.915	< 2e-16 ***
Late2	0.2976810	0.0550675	5.406	6.45e-08 ***
Late3	0.7161429	0.0472809	15.147	< 2e-16 ***

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 31467 on 29654 degrees of freedom
Residual deviance: 26975 on 29647 degrees of freedom
AIC: 26991

Number of Fisher Scoring iterations: 4

```
logreg2 = glm(Default ~ Limit + Late1 + Late2 + Late3, data=loans, family="binomial")  
summary(logreg2)
```

Call:

```
glm(formula = Default ~ Limit + Late1 + Late2 + Late3, family = "binomial",  
    data = loans)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.572950	0.028831	-54.56	< 2e-16 ***
Limit	-0.001878	0.000133	-14.12	< 2e-16 ***
Late1	1.357784	0.041178	32.97	< 2e-16 ***
Late2	0.307697	0.054949	5.60	2.15e-08 ***
Late3	0.718916	0.047186	15.24	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 31467 on 29654 degrees of freedom
Residual deviance: 27042 on 29650 degrees of freedom
AIC: 27052

Number of Fisher Scoring iterations: 4

What do you see in the results? What does the AIC tell you?

Let's find the new model's ROC and AUC next:


```
#1. Make predictions:  
loans$PredProbs2<-predict(logreg2, newdata=loans, type="response")
```

```
#2 Get the ROC and AUC:
roc.pred = prediction(loans$PredProbs2, loans$Default)
# ... which we can then use to actually create the ROC curve
# with the function "performance" (note: we need to store this
# so that we can then draw the curve):
perf = performance(roc.pred, "tpr", "fpr")
#If you don't get what it is doing: ?performance

plot(perf,                                     # the data
      main = "ROC Curve",                     # the chart's title
      xlab = "1 - Specificity",                # the name of the x-axis
      ylab = "Sensitivity",                   # the name of the y-axis
      colorize=TRUE)                          # add color to curve depending on threshold prob.

# ... and add the diagonal corresponding to the Random Assignment
# benchmark model:

abline(0,1) # adds line at intercept 0, with slope 1
```



```
perf_auc = performance(roc.pred, "auc")
as.numeric(perf_auc@y.values)
```

```
[1] 0.7429727
```

How would you characterize the AUC difference between the models?

Which one should we use?

```
MyConfMatrixValues_func(loans$Default, loans$PredProbs2, cutoff=0.55)
```

```
[1] "For a cutoff of 0.55 :"  
[1] "Actual values in the test dataset"  
actual_value  
  0    1  
23045 6610  
[1] "Predicted values in the test dataset"  
predicted_value  
  0    1  
26956 2699  
[1] "There are 22060 True Negatives"  
[1] "There are 4896 False Negatives"  
[1] "There are 985 False Positives"  
[1] "There are 1714 True Positives"  
[1] "As a result:"  
[1] "Accuracy is: 0.8017"  
[1] "True Positive (Sensitivity) Rate is: 0.2593"  
[1] "True Negative (Specificity) Rate is: 0.9573"  
[1] "Precision PPV is: 0.6351"
```

		predicted_value	
actual_value		Pred False	Pred True
Obs False		22060	985
Obs True		4896	1714

Train and Test datasets

So far, we have been training and testing in the same dataset.

Usually, we would train the model on a dataset (or a portion of the data we have) and test it in another dataset (or another portion). These portions are often called “partitions”

```
#Our split should be random but we would also like to have the same "random" results  
#By setting a seed we ensure everyone using the same seed gets the same "random" results.  
  
set.seed(1020, sample.kind = "Rejection")  
df<-loans
```

```
spl = sample(nrow(df), 0.8*nrow(df))
head(spl)
```

```
[1] 1360 7166 1394 13548 368 15061
```

```
# Now lets split our dataset into train and test:
```

```
train.df = df[spl,]
test.df = df[-spl,]
dim(df)
```

```
[1] 29655 11
```

```
dim(train.df)
```

```
[1] 23724 11
```

```
dim(test.df)
```

```
[1] 5931 11
```

With this approach, we would train the model on the train portion of the dataset:

```
logreg2b = glm(Default ~ Limit + Late1 + Late2 + Late3, data=train.df, family="binomial")
summary(logreg2b)
```

Call:

```
glm(formula = Default ~ Limit + Late1 + Late2 + Late3, family = "binomial",
    data = train.df)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.5720841	0.0321928	-48.833	< 2e-16 ***
Limit	-0.0018894	0.0001493	-12.652	< 2e-16 ***
Late1	1.3746490	0.0462526	29.720	< 2e-16 ***
Late2	0.2658272	0.0616894	4.309	1.64e-05 ***
Late3	0.7386307	0.0528218	13.983	< 2e-16 ***

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 25154  on 23723  degrees of freedom
Residual deviance: 21624  on 23719  degrees of freedom
AIC: 21634
```

Number of Fisher Scoring iterations: 4

Then, we can test it with a test dataset:

```
test.df$PredProbs2b<-predict(logreg2b, newdata=test.df, type="response")
MyConfMatrixValues_func(test.df$Default,test.df$PredProbs2b,cutoff=0.55)
```

```
[1] "For a cutoff of 0.55 :"
```

[1] "Actual values in the test dataset"	
actual_value	
0	1
4601	1330

```
[1] "Predicted values in the test dataset"
```

predicted_value	
0	1
5392	539

```
[1] "There are 4402  True Negatives"
[1] "There are 990  False Negatives"
[1] "There are 199  False Positives"
[1] "There are 340  True Positives"
[1] "As a result:"
[1] "Accuracy is: 0.7995"
[1] "True Positive (Sensitivity) Rate is: 0.2556"
[1] "True Negative (Specificity) Rate is: 0.9567"
[1] "Precision PPV is: 0.6308"
```

	predicted_value	
actual_value	Pred False	Pred True
Obs False	4402	199
Obs True	990	340

How do you find the performance to fare between Training and Testing?

Regression analysis are statistically robust, but that will no longer be the case when we move to decision trees, which are more prone to overfitting.